

On The Diffusion of Sticky Particles in 1-D

Joshua DM Hellier* and Graeme J Ackland†

*SUPA, School of Physics and Astronomy, University of Edinburgh,
Mayfield Road, Edinburgh EH9 3JZ, United Kingdom*

(Dated: August 15, 2018)

The 1D Ising model is the simplest Hamiltonian-based model in statistical mechanics. The simplest interacting particle process is the Symmetric Exclusion Process (SEP), a 1D lattice gas of particles that hop symmetrically and cannot overlap. Combining the two gives a model for sticky particle diffusion, SPM, which is described here. SPM dynamics are based on SEP with short-range interaction, allowing flow due to non-equilibrium boundary conditions. We prove that SPM is also a detailed-balance respecting, particle-conserving, Monte Carlo description of the Ising model. Neither the Ising model nor SEP have a phase transition in 1D, but the SPM exhibits a non-equilibrium transition from a diffusing to a blocked state as stickiness increases. We present a fully non-linear, analytic, mean-field solution, which has a crossover from a positive to a negative diffusion constant. Simulations in the positive-diffusion region agree with the analytics. The negative diffusion constant in fact indicates a breakdown of the mean-field approximation, with close to zero flow and breaking into a two-phase mixture, and thus the mean field theory successfully predicts its own demise. The simplicity of the model suggests a wide range of possible applications.


I. INTRODUCTION

Lattice gases are a ubiquitous tool for modeling complex systems from biology to traffic [1–7]. Analytically solvable cases involve non-interacting or excluding particles [8–13], but in any real system of interest the moving objects interact. Many models tackle the situation where the diffusing object interact with the substrate, but despite the clear application-relevance there is surprisingly little work considering interactions between the moving particles themselves. One reason for this is that the interactions introduce nonlinearities in analytical models, which makes them challenging to solve, at least outside of limits in which they can be linearized. This is unfortunate because it is precisely these nonlinearities which introduce interesting behaviors such as discontinuities at the oxide-metal interface or diffusion instability [14, 15].

Another feature of previous models is that the flow is driven by either asymmetry in the dynamics (e.g. ASEP) or an external field which permeates the system (e.g. KLS). In either case, the particles always see a local asymmetry. However in many systems the flow is driven by a pressure or chemical potential difference applied at the boundaries, so that any local asymmetry arises from self-organization. This situation is addressed here.

In this paper we will investigate what we believe to be the simplest such one-dimensional model, the “Sticky Particle Model” or SPM, specified in the top left inset of Fig. 1, which contains such an interaction, and we will explore the impact this has on particle behavior, in particular when observed in the large-scale limit. One might contrast this approach (making a simple microscopic model and trying to learn from it about large-scale

FIG. 1. White circles indicate particles, dark circles indicate empty sites (vacancies). Particles randomly move into adjacent vacancies with rate 1 (having rescaled time for notational convenience), unless there is a particle behind the position they’re moving from, in which case they move with rate λ ; the state of the site next to the position the particle is moving into is irrelevant.



interface growth) with approaches such as the KPZ equation [16–18] (where one analyses the extreme large-scale dynamics using universality classes).

II. THE STICKY PARTICLE MODEL

The SPM is an excluded-particle model in which adjacent particles separate with rate λ and single particles move at rate 1.

It differs from the symmetric exclusion process [19–24]; in that particles “unstuck” with rate λ instead of their

* J.D.M.Hellier@sms.ed.ac.uk

† G.J.Ackland@ed.ac.uk

normal hopping rate, 1. Low λ corresponds to sticky particles, high λ to repelling particles. It could be regarded as a version of the KLS model [25–27] in 1-dimension without an applied field, which is itself similar to the dynamics used to analyze the Ising model by Kawasaki [28]. The KLS model has a field which introduces an asymmetry which drives a flow. It is tacitly assumed that this field is required for the nontrivial flow behavior seen in the model, so the simpler symmetric model has received less attention.

A. Detailed Balance Proof

The SPM is intended to study flow, so it is *defined* by the hopping rates. In Fig. 1 we show all the possible transitions which may occur between local configurations.

Assume that the system is now on a ring, with L lattice sites and N particles. Let us label possible system configurations by ξ and let the number of adjacencies (or “bonds”) between particles be $b(\xi)$. Now for our ansatz, assume that the probability of the system being in state ξ is $\lambda^{-b(\xi)}$. In the top and bottom diagrams of Fig. 1 we can see that the number of bonds on both sides is the same, as are the transition rates back and forth; thus our ansatz holds for these states, as it predicts the probabilities of the left and right configurations are the same. The middle two diagrams are equivalent; in the upper diagram a bond is formed going left to right and then broken going right to left, so the probability of being in the left state is λ times that of being in the right state. This is again in agreement with the detailed balance criterion. As these are the only types of transition that may occur on a ring, we have proven that the closed system obeys detailed balance, with an energy proportional to $b(\xi) \ln \lambda$.

B. Equivalence to Ising Model at Equilibrium

Since the model obeys detailed balance, there must be an associated Hamiltonian. This is simply proportional to the number of particles stuck together:

$$H = q \sum_i p_i p_{i+1} = qb(\xi) \quad (1)$$

where p_i is 1 for a particle or zero for an empty site. A simple transformation $s_i = 2p_i - 1$ shows this to be equivalent to the Ising Hamiltonian

$$H = q \sum_i \left[\frac{1}{4} s_i s_{i+1} + \left(\frac{1}{2} s_i + \frac{1}{4} \right) \right] \quad (2)$$

The contribution from the final term is constant, as only hopping moves are allowed so $\sum s_i$ is constant. So at equilibrium the SPM samples the 1D Ising model with fixed magnetization, a fact which was used to validate the

codes. Compatibility between the detailed balance condition, the standard Boltzmann distribution at temperature T and our defined rates requires that $q = k_B T \ln \lambda$.

C. Flow Rate Calculations

The SPM, just like the KLS model and ASEP, can exhibit flow. In this Letter we will be principally concerned with the computation of the flow that results when an SPM system of L sites is connected to (unequal) particle reservoirs at either end. It is actually possible to “analytically” solve the SPM on a finite domain with fixed boundary densities by analyzing the transition rate matrix which represents that system. Specifically, we build a sparse representation of the matrix, and then use the Python routine `scipy.sparse.linalg.eigs` to find the eigenvector corresponding to the eigenvalue closest to zero (which we expect to be exactly zero in the absence of numerical uncertainty). We can then retrieve the current through the system in a steady state by projecting that eigenvector using a current operator. Unfortunately the dimensionality of this matrix scales exponentially with the system size, which imposes upon us a memory limit restricting us to relatively small systems. In spite of this, we have managed to compute properties such as the steady-state flow rate for systems with sizes as large as 12 lattice sites, which are large enough to produce consistent outputs comparable to our other methods, whilst being computationally cheap enough to compute en masse. In addition, we have also performed Monte Carlo simulations using a couple of different methods, which allows us to better probe the large-system regime.

In the simulations and “analytics”, we can vary ρ_0 , ρ_L and λ to address three questions: how does λ affect the flow, how does the particle current depend on the driving force $(\rho_0 - \rho_L)/L$, and how does the density of particles in the system vary?

III. SIMULATIONS

In the SPM the bulk dynamics are perfectly symmetric, and flow is driven by setting ρ at the boundaries, with the system self-organizing to give the steady-state density. We hypothesise that the flow rate, Hamiltonian energy and particle density of the system are fully determined by (ρ_0, ρ_L) and the unsticking rate λ ; this has been so far confirmed in our “analytics” by the existence and uniqueness of the steady flow state in every instance we have tried.

We can simulate flow in a bounded domain using Monte Carlo methods. In the bulk, the transition rates are simply those described in Fig. 1. At the boundaries there are 2 layers of lattice sites that switch between being full and empty with rates such that the time-averaged occupation (ρ_0, ρ_L) defines the boundary conditions; the system is therefore open, and particles appear and disap-

pear at boundary layers. These boundaries are equivalent to having particle reservoirs attached to the edges of the domain.

We chose to calculate primarily by using the KMCLib[29] package, which implements the Kinetic Monte Carlo algorithm (essentially the same as the Gillespie algorithm[30–32]) on lattice systems. The codes used are kept here [33]. Independent of the KMCLib code, we wrote a simple Metropolis-Hastings algorithm [34] which randomly selects single particle hops; this is more efficient than using KMC for values of λ which are relatively close to 1 as the algorithm is simple and the acceptance rate is relatively high, so we can generate better statistics using this method in that regime. We calculate the flow from the number of particles entering and leaving the system at the boundaries. Since the model is defined in terms of *rates*, the flow in “particles per unit time” is a well defined quantity.




FIG. 2. The variation of (normalized) flow rate with λ in the SPM with fixed boundary densities, on both logarithmic and linear scales. Crosses indicate results from the KMC simulations, triangles those from the traditional Metropolis-Hastings simulations and the joined circles the current observed by finding the least dominant eigenvector in the transition rate matrix. The dashed lines are the corresponding MFT prediction for the current, and color indicates the boundary conditions used. The flow rate is normalized in the sense that it has been multiplied by the system size, so that we can easily compare between different system sizes. We do have estimates of the standard error in the flow rate, but as we have an abundance of data points we feel that the spread of the points on the graph illustrates it well enough.

Using KMCLib we studied systems of length 64 (lengths

32, 128 and 256 give similar results as shown in Fig. 5.), running them for 400000 Gillespie steps for equilibration followed by 10000 measurement runs of 1000 steps interspersed with relaxation runs of 16000 steps. This way we could gather statistics about flow rates and densities in a well-equilibrated system. Specifically, we generate a pool of 10000 samples of flow rate and density, from which we can calculate estimates of the descriptive statistics of both quantities. These calculations, in which we held the boundary densities constant whilst varying λ , were repeated using Metropolis-Hastings for a length 100 system, and we have transition rate matrix currents for a system of length 10. These are plotted on both logarithmic and linear axes in Fig. 2. respectively.




FIG. 3. As the logarithmic plot in Fig. 2, but over a much wider range of λ values. Notice how the data series from the two Monte-Carlo methods broaden and flatten as λ becomes very small; this is because of poor statistics due to the low flow rate.

A. Diffusion Coefficient


We compare an MFT prediction (see IV) and the KMC numerical results for the diffusion constant in Fig. 11. We see that MFT and simulation agree well for low stickiness, and both show the symmetry about $\rho_M = \frac{2}{3}$. For high stickiness, where the MFT prediction gives negative diffusion constant, we actually see low positive values for the current. It should be noted that the MFT assumes that $\rho = \rho_M$ throughout, whereas in the simulations ρ tends to be much higher.

B. Particle Motion During Flow

It is instructive to get an overview of how the particles move during flow. Fig. 4 show a plot of the flow structure in the slow-flow regime. In very short time averages the “striped” pattern indicates separation into dense and sparse regions, with an overall concentration gradient arising from the relative size of such regions. Particle/vacancy diffusion through the empty/full regions can

be seen. As the averages are taken over longer times the blocks themselves appear to diffuse. Averaged over the entire simulation (not shown) the averaging simply gives a smooth density gradient.

FIG. 4. Indicative spacetime flow pattern for sticky free-flow [$\lambda = \frac{3}{20}, (\rho_0, \rho_L) = (\frac{3}{4}, \frac{1}{4})$]; other combinations shown in the supplementary materials. Time runs along the x-axis, space (1 pixel=1 site) along the y-axis, with grayscale tone (black being empty, white being full) illustrating average site occupation over (clockwise from top left) $\frac{1}{32}$, 1, 8 and 32 Gillespie steps per site respectively.



Interesting structure is visible, the dynamics appearing as a random walk with some tendency for particles to clump; over longer timescales the diffusive behavior is more evident, with a textured structure suggesting characteristic velocity of particles or vacancies through emergent correlated clumps. Additional plots can be found in the supplementary materials.

C. Transition in flow character

Fig. 3 shows how the current varies with stickiness at fixed driving for a very broad range of λ . At high λ the rate is simply proportional to λ for any forcing. This result is far from trivial - it means that the overall flow is determined by the *faster* rate (λ), not the slower rate (1). That the MFT averages over the two is unsurprising (Eq. 8), but for the simulation to avoid having a “rate limiting step” requires the system to fill to sufficient density that there are always particles in contact to repel one another.




FIG. 5. The system size dependence of the normalized flow rate as a function of λ , computed using the KMC method. These calculations all have boundary conditions (0.75, 0.25), and the flow rate has been normalized via multiplication by the system size (displayed on legend). Size 64 systems are a good compromise in terms of size and accuracy for a given amount of computational time.

When $\lambda = \lambda_c \sim 0.2$ the simulations show a transition to a different behavior, with higher or lower λ_c depending on the boundary conditions and the strength of the driving force. The current and its fluctuation remains finite, but there is a distinctive peak in the particle density fluctuation $\langle (\rho - \bar{\rho})^2 \rangle$ (Fig. 6). The width of the peak is independent of the system size which suggests a continuous phase transition from the free flowing to the “stuck” regime.




FIG. 6. The scaling of the normalized fluctuation in the overall system density with system size, as a function of λ and with boundary conditions (0.6, 0.4). This was computed using the traditional Metropolis-Hastings method. The normalization was achieved by multiplying by the system size, L . The great deal of coherence between these results regardless of size suggests that the fluctuation is of the normal Gaussian type, and not the kind of anomalously-scaling fluctuation associated with equilibrium phase transitions.

At very low λ it becomes exceptionally rare for a particle or vacancy to traverse the whole system, so the two Monte-Carlo methods are unable to generate good statistics, and the flow measurement is dominated by the noise

of boundary fluctuations. Another way of looking at this is that these methods begin failing to equilibrate properly in the time they are provided with provided, on the grounds that the equilibration time varies as $\mathcal{O}(\lambda^{-1})$ as $\lambda \rightarrow 0$ (see Fig. 7). However, there is a regime when $0.04 \leq \lambda \leq 0.1$ where the mean flow clearly displays power-law behavior, with $\mathcal{O}(\lambda^k)$ with $k \sim 4$; between this region and the λ^1 region there is a bend as we switch between the power laws. This bend, with its accompanying peak in the density fluctuation, is the “smoking gun” of the nonequilibrium transition we claim to have found.




FIG. 7. The moduli of the 5 smallest eigenvalues of the transition rate matrix for the SPM for a system of size 10 as a function of λ , computed using sparse numerical linear algebra. They are all negative real to within numerical precision, except the very smallest, which is numerically 0 (and therefore not displayed). Notice the general trend, whereby they scale roughly linearly with λ . As these eigenvalues correspond to the decay rates of the dominant unsteady states in the system, this implies that the time taken for a simulated system to relax to its equilibrium state scales inversely proportional with λ , which helps explain the poor convergence of our Monte Carlo simulations in that regime. Observe also that for intermediate λ eigenvalues often merge or cross; this implies the existence of rich behavior, and suggests that finding a full analytic solution to for this model might be difficult.

D. Effective diffusion constant

Using boundary conditions $(\rho_0, \rho_L) = (\rho_M + \frac{1}{2}\delta\rho, \rho_M - \frac{1}{2}\delta\rho)$ we demonstrate the dependence of current on boundary density difference and λ (Fig. 10). This shows that the transition to the stuck phase is suppressed by stronger driving forces (large $|\delta\rho|$)

We use the limit of small $\delta\rho$ to calculate the effective diffusion constant (Fig. 11). This is normalized to 1 in the case of $\lambda = 1$, which is just SEP. The $\rho_M = 0$ limit corresponds to free flow of particles, so the diffusion constant here does to 1. Similarly the $\rho_M = 1$ limit is flow of vacancies, so $D \rightarrow \lambda$. One might expect a monotonic variation between these limits, but surprisingly the simulations show that there is always an extremal value for D close to $\rho_M = 2/3$: this is a minimum for $\lambda < 1$ and a maximum for $\lambda > 1$

E. Self-organized density

We calculated the time-averaged total number of particles in the system by updating a histogram of particle numbers as the simulation progresses. In each of our calculations, we make the initial configuration by randomly filling the system with particles and vacancies in such a way that the initial density should be $\frac{1}{2}(\rho_0 + \rho_L)$, and then run the system for a sufficient number of equilibration steps to destroy any initial transients.

In SEP, ($\lambda = 1$) the density varies linearly across the system from ρ_0 to ρ_L , as one expects for a diffusion process. However, for sticky particles, $\lambda < 1$, the density rises sharply near to the boundary, then has a linear profile about a value higher than ρ_M (Fig. 8). One might view this as particles being sucked into the system to lower their Hamiltonian Energy, but such a notion can be dismissed since the internal density *also rises for repelling particles*, $\lambda > 1$. The asymptotic values for the bulk density at high and low λ appear to be 1 and $\frac{2}{3}$ respectively, regardless of the mean boundary density ρ_M (see Fig. 9).

F. Review of the simulations

Simulation of the sticky particle model reveals a number of curious and unexpected features which have no equivalent in either SEP or the 1D Ising model.

- A transition in the dependence of the current on λ from $\propto \lambda$ to $\propto \lambda^4$.
- The mean density of the system has complicated variation with λ .
- The effective diffusion constant has an extremal value at intermediate boundary density.

To get further understanding of these behaviors, we tackle the system analytically, using mean field theory and transition rate matrices.

IV. MEAN FIELD THEORY FOR FLOW

Let the spacing between lattice sites be a , let τ_0 be the free-particle hopping timescale, and the time-averaged (or ensemble-averaged, assuming ergodicity) occupation probability of the i^{th} lattice site be ρ_i . We will introduce $\zeta = 1 - \lambda$ here for convenience.

A. Lattice Mean Field Theory

Let the ensemble-averaged occupation probability of the i^{th} site at time t be $\rho_i(t)$. In the mean-field approximation this is assumed to be independent of $\rho_j(t)$ for

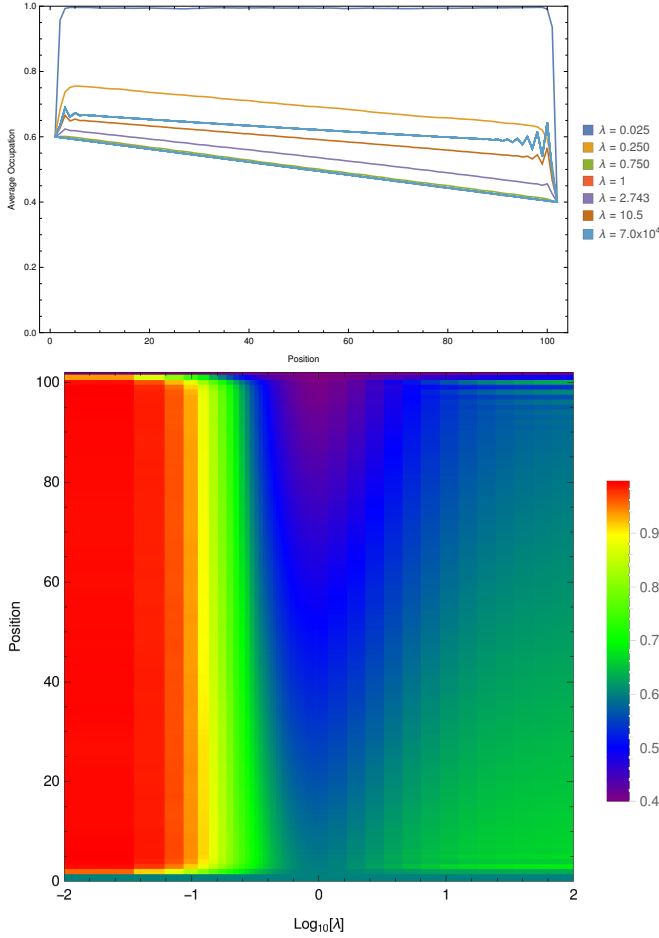


FIG. 8. The variation of the time-averaged density profile with respect to λ . The top panel shows this profile for a selection of values for λ in a system of size 100, with boundary conditions (0.6, 0.4); the bottom panel displays this variation in the form of a colored density plot. Again, we have used the traditional Metropolis-Hastings method for these calculations. The oscillatory behavior near the boundary is a robust feature of the profile at high λ .

$j \neq i$ at equal times. Therefore, if the i^{th} site is occupied, then the rate at which it empties is the sum of contributions from the four permutations of the $(i+1)^{\text{th}}$ and $(i-1)^{\text{th}}$ occupancy:

$$\begin{aligned} & \frac{1}{\tau_0} (1 - \rho_{i-1}) [(1 - \rho_{i+1}) + \lambda \rho_{i+1}] \\ & + \frac{1}{\tau_0} (1 - \rho_{i+1}) [(1 - \rho_{i-1}) + \lambda \rho_{i-1}]. \end{aligned} \quad (3)$$

Similarly, if the $(i)^{\text{th}}$ site is unoccupied, it fills with rate which depends on the occupation probability of the neighbours, and whether they are stuck to the $(i+2)^{\text{th}}$ and $(i-2)^{\text{th}}$ sites respectively. Therefore, an unoccupied




FIG. 9. The overall average density of a system of size 64, computed using KMC, with boundary conditions as specified in the plot. Notice how the density converges to 1 and $\sim \frac{2}{3}$ for extreme low and high λ respectively.

i^{th} site fills with rate

$$\frac{1}{\tau_0} \{ \rho_{i+1} [\lambda \rho_{i+2} + (1 - \rho_{i+2})] + \rho_{i-1} [\lambda \rho_{i-2} + (1 - \rho_{i-2})] \}. \quad (4)$$

If we now multiply the filling/emptying rates of site i by the probability of being empty/full respectively, we obtain the final equation for the site occupation

$$\begin{aligned} \tau_0 \frac{\partial \rho_i}{\partial t} = & (1 - \rho_i) [(1 - \zeta \rho_{i-2}) \rho_{i-1} + (1 - \zeta \rho_{i+2}) \rho_{i+1}] \\ & - \rho_i [2\zeta \rho_{i-1} \rho_{i+1} - (3 - \zeta) (\rho_{i-1} + \rho_{i+1}) + 2]. \end{aligned} \quad (5)$$

B. Diffusion Equation: MFT Continuum Limit

To obtain the continuum limit of the MFT we substitute $\rho_i(t) \rightarrow \rho(x, t)$, $\rho_{i+m}(t) \rightarrow \rho(x + am, t)$ into Eq. 5. Then, using a Taylor expansion around x for small a , neglecting terms of $\mathcal{O}(a^4)$, and collecting terms we find that

$$\begin{aligned} \tau_0 \frac{\partial \rho}{\partial t} = & a^2 [1 - \zeta \rho (4 - 3\rho)] \frac{\partial^2 \rho}{\partial x^2} \\ & + 2a^2 \zeta (3\rho - 2) \left(\frac{\partial \rho}{\partial x} \right)^2 + \mathcal{O}(a^4), \end{aligned} \quad (6)$$

which can be factorized into the more familiar form of a continuity equation

$$\frac{\partial \rho}{\partial t} = \frac{a^2}{\tau_0} \frac{\partial}{\partial x} \left\{ [1 - \zeta \rho (4 - 3\rho)] \frac{\partial \rho}{\partial x} \right\}, \quad (7)$$

having dropped the higher-order terms. From this we can identify the flow $J(x)$ as the term in the curly brackets, and define a mean field diffusion constant (see Fig.11):

$$D_{MFT}(\zeta, \rho) = \frac{a^2}{\tau_0} [1 - \zeta \rho (4 - 3\rho)] \quad (8)$$

We note that the MFT diffusion constant can become negative. The density at which this occurs is given by finding the roots ρ_c of the RHS of Eq. 8:

$$\rho_c = \frac{2}{3} \pm \frac{1}{3} \sqrt{4 - 3/\zeta}. \quad (9)$$

The possible values for ρ_c are only real (and therefore physically possible) for $\zeta < 0.75$. So in the limit of very sticky particles ($\lambda < 0.25$) the MFT predicts a transition between forward and backwards diffusion at some densities.

C. Limiting Cases

In order to understand the implications of the MFT, let us consider some limits. As $\zeta \rightarrow 0$ (i.e. as the model becomes a simple exclusion model), $D \rightarrow \frac{a^2}{\tau_0}$. Likewise, in the dilute limit $\rho \rightarrow 0$, $D \rightarrow \frac{a^2}{\tau_0}$, reflecting the fact that it becomes a dilute lattice gas and therefore the interactions between particles become irrelevant as they never meet. Conversely, in the full limit $\rho \rightarrow 1$, $D \rightarrow \frac{\lambda a^2}{\tau_0}$; this is because we now have a dilute gas of vacancies, which hop with rate $\frac{\lambda}{\tau_0}$.

One may observe that the MFT has a symmetry under $\rho \mapsto \frac{4}{3} - \rho$; thus, the dynamics should be symmetric under a density profile reflection around $\rho = \frac{2}{3}$. This is where D always attains its extremal value, $\frac{a^2}{\tau_0} [1 - \frac{4}{3}\zeta]$, hence for $\zeta > 3/4$ the diffusion coefficient becomes negative in regions with $\frac{2}{3} - \frac{\sqrt{\zeta(4\zeta-3)}}{3\zeta} < \rho < \frac{2}{3} + \frac{\sqrt{\zeta(4\zeta-3)}}{3\zeta}$.

Finally, it is possible to show that solutions to the continuum MFT containing domains with a negative diffusion coefficient are linearly unstable; thus, if we try to have a flow containing ρ for which $D(\rho) < 0$, the density of the medium should gravitate towards one of the two densities for which $D(\rho) \sim 0$. Instead of observing “backwards diffusion” we would see an extremely slow flow or no flow at all.

D. Steady State Flow

It is possible to solve the continuum MFT in a steady state on a finite domain, say $x \in (0, L)$. Steady state implies that there is no build-up of particles $\frac{\partial \rho}{\partial t} = 0$, so the flow is constant through the system.

Using the constant-flow requirement for continuity $J(x) = J_0 = \text{const.}$, and by integrating both sides of Eq. 7 with respect to x we find that

$$\int_{x_0}^x J(x') dx' = -\frac{a^2}{\tau_0} \int d\rho [1 - \zeta \rho (4 - 3\rho)] \quad (10)$$

$$(x - x_0) J_0 = -\frac{a^2}{\tau_0} [\rho + \zeta(\rho - 2)\rho^2]. \quad (11)$$

The density profile across the system is given by solving for $\rho(x)$. Since Eq.11 is cubic, the solution for density

$\rho(x)$ is non-unique for cases of high ζ . Thus in the limit of high stickiness, the MFT is unable to make a unique prediction for the density. Furthermore, except in the SEP case $\zeta = 0$, the density will not vary linearly across the system.

E. Dirichlet Boundary Conditions

The constants x_0 and J_0 are the boundary conditions for the flow, but they need not correspond to a physically-realizable situation. For driven systems it is more convenient to consider fixing the density at each end.

If we impose such Dirichlet boundary conditions on this system, say $\rho(0) = \rho_0$ and $\rho(L) = \rho_L$, we find that

$$J_0 = \frac{a^2}{L\tau_0} [\rho_0 - \rho_L + \zeta (\rho_0 [\rho_0^2 - 2] - \rho_L [\rho_L^2 - 2])]. \quad (12)$$

This equation can be used for direct comparison with the simulations (Fig. 10 and Fig. 11). In general, the agreement is good, except for the region where the MFT predicts negative flow.

F. Interpretation of MFT

The mean field theory enables us to predict flow behavior of the MFT. It recovers well-known limiting cases such as SEP ($\lambda = 0$), however at high stickiness ($\lambda < 0.25$) it predicts its own demise, with unphysical negative diffusion constants and by having multiple solutions for the density at some positions in the system, meaning that the MFT does not give a unique density profile. This breakdown of the MFT corresponds to the transition to slow flow observed in the simulation.

In some conditions MFT predicts densities greater than 1. One might guess that when the MFT offers two possible values for the density, it will correspond to a phase separation in the actual system. Furthermore, the MFT prediction of a maximum diffusion constant at a density of $\rho = 2/3$ suggests that this value of ρ might be favoured for strongly driven flows.

A curious feature of the MFT diffusion (Fig. 11) is that for fixed stickiness there are *two possible densities* giving the same diffusion constant. Thus it is possible to have a steady state flow with phase separation into regions of high and low density. This echoes the situation seen in short time-averages of the simulation (Fig.4) where blocks of high and low density are evident, as well as the oscillatory behavior near the boundaries at high λ (Fig. 8).

Regarding the phase transition, we should note that we only have an MFT prediction for the flow rate as a function of λ , since $\rho(x)$ stops being unique when λ drops below $\frac{1}{4}$, and so the MFT lacks predictive power. For low-stickiness, when $\lambda > \frac{1}{4}$, the MFT is in very good agreement with the simulations, and this continues as λ

stretches into the thousands, where the mean flow varies as $\mathcal{O}(\lambda^1)$.

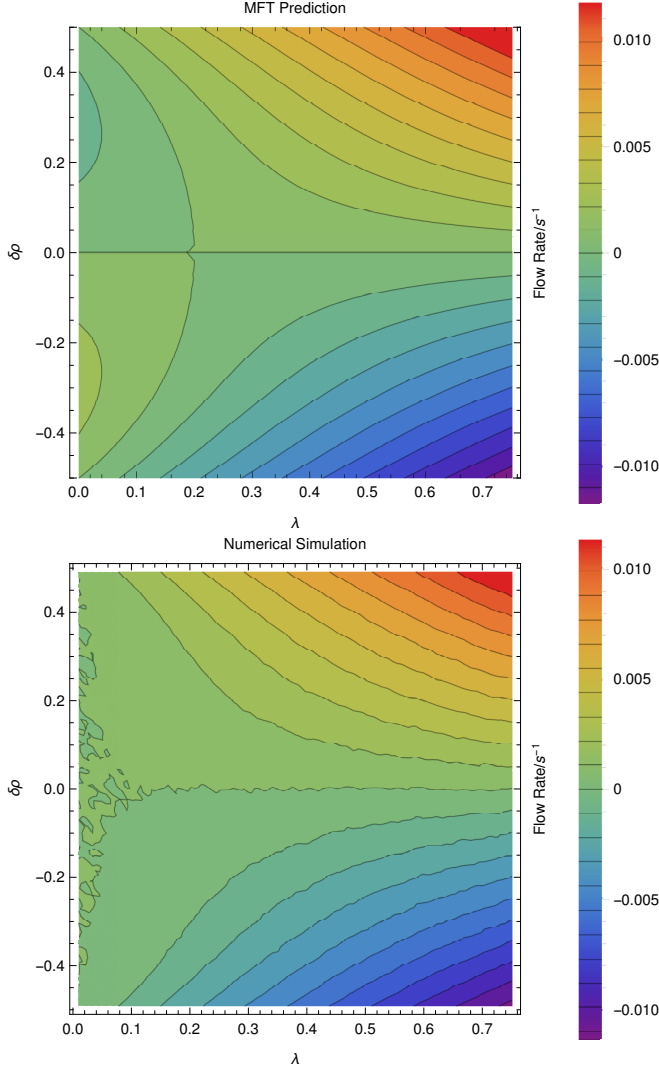


FIG. 10. Flow rate mean observed when varying the difference $\delta\rho$ between the boundary concentrations $(\rho_0, \rho_L) = (\rho_M + \frac{1}{2}\delta\rho, \rho_M - \frac{1}{2}\delta\rho)$ and λ (The top panel is the MFT prediction for the flow rate, whilst bottom shows the observed mean flow rate). We chose $\rho_M = \frac{1}{2}$, as this gives us the biggest range of $\delta\rho$ to investigate. These calculations were performed with the same run parameters (system length etc) as described in Fig. 11. The MFT prediction (top) shows a region of negative flow - this occurs at $\zeta > \frac{4}{5}$ ($\lambda < 0.2$) for the weakest driving force, with additional stickiness required when the driving force increases. The simulation exhibits close to zero flow in that region (bottom). Away from this region, the MFT and the simulation are in good agreement.

The MFT prediction for the mean flow is again a good fit until λ becomes sufficiently small, when the simulation show no evidence of negative diffusion; rather the flow becomes critically slow for very sticky particles. The higher moments of the simulated flow (e.g. variance) do not show peaks, indicating that hard transitions are not

occurring. Finally, the simulated density is very close to the average of the boundary densities until λ drops below $1/4$, at which point the system fills.

Continuing to specify the boundary densities to be $(\rho_0, \rho_L) = (\rho_M + \frac{1}{2}\delta\rho, \rho_M - \frac{1}{2}\delta\rho)$ for some given ρ_M , we can keep $\delta\rho$ relatively small, so that J varies approximately linearly with $\delta\rho$; thus if we calculate J for a series of small $\delta\rho$, we can perform linear regression to find $D_{\text{Eff}} = \left. \frac{\partial J}{\partial \delta\rho} \right|_{\delta\rho=0}$, the effective diffusion coefficient. Computing this for different (ρ_M, λ) combinations yields results that can be compared with Eq. 12.

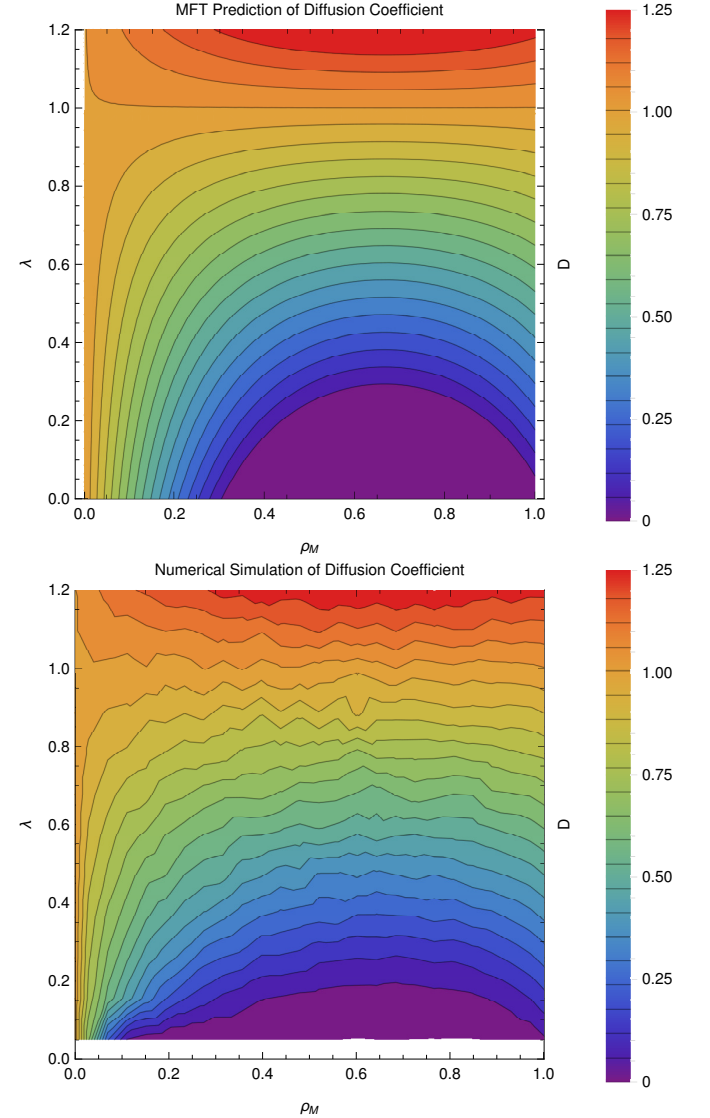


FIG. 11. Comparison of effective diffusion coefficient D in the MFT (top) and in direct simulation (bottom) as a function of density and stickiness. The white region is where the MFT gives negative diffusion. The simulations used 124 sites averaged over $\sim 10^9$ steps at each of $12 \times 24 \times 16$ ($\lambda, \rho_M, \delta\rho$) combinations. Full details in the supplementary materials.

